

Energy-based Models

--Boltzmann Machine

Hao Dong

Peking University


Content



- Boltzmann Machine
 - Introduction
 - Training without hidden neurons
 - Training with hidden neurons
 - Summary
- Restricted Boltzmann Machine
- Deep Boltzmann Machine

- Boltzmann Machine
 - Introduction
 - Training without hidden neurons
 - Training with hidden neurons
 - Summary
- Restricted Boltzmann Machine
- Deep Boltzmann Machine

Recap: The Stochastic Hopfield Net



$$z_i = \frac{1}{T} \sum_{j \neq i} w_{ij} y_j + b_i$$

$$E(Y) = \sum_{i < j} -w_{ij} y_i y_j - b_i y_i$$

$$P(y_i = 1) = \sigma(z_i)$$

$$P(y_i = -1) = 1 - \sigma(z_i)$$

$$P(Y) = \frac{\exp(-E(Y))}{\sum_{s'} \exp(-E(Y'))}$$

- The stochastic Hopfield net models a probability distribution over states
 - The state Y is a binary sequence
 - It models a Boltzmann distribution
- The probability that the network will be in any state is $P(Y)$
 - Generative model: generates states according to $P(Y)$

Recap: The Stochastic Hopfield Net



- $P(Y) = P(y_i = 1 | y_{j \neq i}) P(y_{j \neq i})$
- Consider two states Y and Y' with the i -th bit in the $+1$ and -1
- $$\begin{aligned} \log P(Y) - \log P(Y') &= \log P(y_i = 1 | y_{j \neq i}) - \log P(y_i = -1 | y_{j \neq i}) \\ &= \log \frac{P(y_i = 1 | y_{j \neq i})}{1 - P(y_i = 1 | y_{j \neq i})} \end{aligned}$$

Recap: The Stochastic Hopfield Net

- Consider two states Y and Y' with the i -th bit in the +1 and -1



$$P(Y) = \frac{\exp(-E(Y))}{\sum_{s'} \exp(-E(Y'))}$$

- $\log P(Y) = -E(Y) + C$
- $E(Y) = -\frac{1}{2} (E_{without\ i} + \sum_{j \neq i} w_{ij} y_j + b_i)$
- $E(Y') = -\frac{1}{2} (E_{without\ i} - \sum_{j \neq i} w_{ij} y_j - b_i)$
- $$\begin{aligned} \log P(Y) - \log P(Y') &= E(Y') - E(Y) \\ &= \sum_{j \neq i} w_{ij} y_j + b_i \end{aligned}$$




Recap: The Stochastic Hopfield Net

- $\log \frac{P(y_i = 1 | y_{j \neq i})}{1 - P(y_i = 1 | y_{j \neq i})} = \sum_{j \neq i} w_{ij} y_j + b_i$
- We get:
 - $P(y_i = 1 | y_{j \neq i}) = \frac{1}{1 + e^{-(\sum_{j \neq i} w_{ij} s_j + b_i)}}$
- It's a logistic!



Recap: The Stochastic Hopfield Net


- We can make Hopfield net stochastic
 - Each neuron responds probabilistically
 - More in accord with Thermodynamic models
 - More likely to escape spurious “weak” memories



$$z_i = \frac{1}{T} \sum_{j \neq i} w_{ij} y_j + b_i$$
$$P(y_i = 1) = \sigma(z_i)$$
$$P(y_i = -1) = 1 - \sigma(z_i)$$

Running the network


- Initialise the neurons
- Cycle through the neurons and set the neurons to 1/-1 according to the probability
- Until convergence, sample the individual neurons



$$z_i = \frac{1}{T} \sum_{j \neq i} w_{ij} y_j + b_i$$
$$P(y_i = 1) = \sigma(z_i)$$
$$P(y_i = -1) = 1 - \sigma(z_i)$$

The overall probability


- The probability of any state y can be shown to be given by the Boltzmann distribution
 - $E(y) = -\frac{1}{2} y^T W y$
 - $P(y) = C \exp(-\frac{E(y)}{T})$
- Minimising energy maximises log likelihood
- The parameter of the distribution is the weights matrix W



$$z_i = \frac{1}{T} \sum_{j \neq i} w_{ij} y_j + b_i$$
$$P(y_i = 1 | y_{j \neq i}) = \sigma(z_i)$$
$$P(y_i = -1 | y_{j \neq i}) = 1 - \sigma(z_i)$$

The overall probability


- The probability of any state y can be shown to be given by the Boltzmann distribution
 - $E(y) = -\frac{1}{2} y^T W y$
 - $P(y) = C \exp(-\frac{E(y)}{T})$
- The conditional distribution of individual bits in the sequence is a logistic
- We call this **Boltzmann Machine**



$$z_i = \frac{1}{T} \sum_{j \neq i} w_{ij} y_j + b_i$$
$$P(y_i = 1 | y_{j \neq i}) = \sigma(z_i)$$
$$P(y_i = -1 | y_{j \neq i}) = 1 - \sigma(z_i)$$

Boltzmann Machine


- It can be viewed as a generative model
- Probability of producing any binary vector y :
 - $E(y) = -\frac{1}{2}y^T W y$
 - $P(y) = C \exp(-\frac{E(y)}{T})$



$$z_i = \frac{1}{T} \sum_{j \neq i} w_{ij} y_j + b_i$$
$$P(y_i = 1 | y_{j \neq i}) = \sigma(z_i)$$
$$P(y_i = -1 | y_{j \neq i}) = 1 - \sigma(z_i)$$


The capacity of Boltzmann Machine

- The network can store up to N N -bit patterns
- How to increase the capacity?




The capacity of Boltzmann Machine

- Add some nodes
 - We don't care the value of these nodes
 - Only serve to increase the capacity
 - Termed **Hidden Neurons**
- The neurons whose values are important: **Visible Neurons**



Hopfield Net v.s. Boltzmann machine



$$E(Y) = \sum_{i < j} -w_{ij} y_i y_j - b_i y_i$$
$$P(Y) = \frac{\exp(-E(Y))}{\sum_{s'} \exp(-E(Y'))}$$


- Hopfield net
 - Learn weights to “remember” target states and “dislike” other states
 - State: binary pattern of all the neurons
- Boltzmann machine
 - Learn weights to assign more probability to patterns we “like” and less to other patterns

Content



- Boltzmann Machine
 - Introduction
 - Training without hidden neurons
 - Training with hidden neurons
 - Summary
- Restricted Boltzmann Machine
- Deep Boltzmann Machine

Training the network



$$E(Y) = \sum_{i < j} -w_{ij} y_i y_j - b_i y_i$$
$$P(Y) = \frac{\exp(-E(Y))}{\sum_{s'} \exp(-E(Y'))}$$

- First we consider the setting without hidden neurons
- Boltzmann machine
 - Given a set of training inputs Y_1, Y_2, \dots, Y_N
 - Assign higher probability to patterns seen more frequently
 - Assign lower probability to patterns that are not seen at all

Maximum likelihood training

$$\log(P(Y)) = \left(\sum_{i < j} w_{ij} y_i y_j \right) - \log \left(\sum_{Y'} \exp \left(\sum_{i < j} w_{ij} y'_i y'_j \right) \right)$$

$$\mathcal{L} = \frac{1}{N} \sum_{Y \in S} \log(P(Y))$$

$$= \frac{1}{N} \sum_Y \left(\sum_{i < j} w_{ij} y_i y_j \right) - \log \left(\sum_{Y'} \exp \left(\sum_{i < j} w_{ij} y'_i y'_j \right) \right)$$

- The loss function is average log likelihood of training vectors $S = \{Y_1, Y_2, \dots, Y_N\}$
 - should be maximised
 - In the first summation, y_i and y_j are bits of Y
 - In the second summation, y'_i and y'_j are bits of Y' (vectors outside S)

Maximum likelihood training

$$\mathcal{L} = \frac{1}{N} \sum_Y \left(\sum_{i < j} w_{ij} y_i y_j \right) - \log \left(\sum_{Y'} \exp \left(\sum_{i < j} w_{ij} y'_i y'_j \right) \right)$$
$$\frac{d\mathcal{L}}{dw_{ij}} = \frac{1}{N} \sum_Y y_i y_j \quad -?$$

- Use gradient ascent
- The first term is easy to calculate
 - The average $y_i y_j$ over all training vectors
- But the second term is the sum of almost all states
 - exponential number!

The second term

$$\begin{aligned}\frac{d \log(\sum_{Y'} \exp(\sum_{i < j} w_{ij} y'_i y'_j))}{dw_{ij}} &= \sum_{Y'} \frac{\exp(\sum_{i < j} w_{ij} y'_i y'_j)}{\sum_{Y''} \exp(\sum_{i < j} w_{ij} y''_i y''_j)} y'_i y'_j \\ &= \sum_{Y'} P(Y') y'_i y'_j\end{aligned}$$

- The second term is the expected value of y'_i, y'_j over all possible values of the state
- We cannot compute it exhaustively, then how?
- Sampling!

The second term

$$\begin{aligned}\frac{d \log(\sum_{Y'} \exp(\sum_{i < j} w_{ij} y'_i y'_j))}{dw_{ij}} &= \sum_{Y'} P(Y') y'_i y'_j \\ &= \frac{1}{M} \sum_{Y' \in Y_{\text{sample}}} y'_i y'_j\end{aligned}$$

- The expectation can be estimated as the average of samples drawn from the distribution
- How to sample?

Gibbs Sampling


- A special Metropolis-Hastings algorithm
- Use the conditional distribution
- Suppose y_1, y_2, \dots, y_n :
 - Randomly set values to them
 - Update y_i based on $P(y_i | y_{j \neq i})$
 - Get a Markov Chain
 - Skip the first several samples and sample at intervals
- The samples are approximately close to the joint distribution

Maximum Likelihood Training

$$\frac{d\mathcal{L}}{dw_{ij}} = \frac{1}{N} \sum_Y y_i y_j - \frac{1}{M} \sum_{Y' \in Y_{sample}} y'_i y'_j$$
$$w_{ij} = w_{ij} + \alpha \frac{d\mathcal{L}}{dw_{ij}}$$

- The overall gradient ascent rule


Training Process



$$\frac{d\mathcal{L}}{dw_{ij}} = \frac{1}{N} \sum_Y y_i y_j - \frac{1}{M} \sum_{Y' \in Y_{sample}} y'_i y'_j$$
$$w_{ij} = w_{ij} + \alpha \frac{d\mathcal{L}}{dw_{ij}}$$

- Initialise weights
- Obtain “state samples”
- Compute gradient and update weights
- Iterate until convergence


Training Process



$$\frac{d\mathcal{L}}{dw_{ij}} = \frac{1}{N} \sum_Y y_i y_j - \frac{1}{M} \sum_{Y' \in Y_{sample}} y'_i y'_j$$

$$w_{ij} = w_{ij} + \alpha \frac{d\mathcal{L}}{dw_{ij}}$$

- Similar to the update rule for Hopfield network




Content




- Boltzmann Machine
 - Introduction
 - Training without hidden neurons
 - Training with hidden neurons
 - Summary
- Restricted Boltzmann Machine
- Deep Boltzmann Machine

Training with hidden neurons

- For a given pattern of visible neurons, there are many hidden patterns (2^K)
- We want to choose the one with lowest energy
 - But exponential search space is exponential!



Training the network




$$P(Y) = \frac{\exp(-E(Y))}{\sum_{s'} \exp(-E(Y'))}$$

$$P(Y) = P(V, H)$$

$$P(V) = \sum_H P(Y)$$

- $Y=(V, H)$
 - V: output of the visible neurons
 - H: output of the hidden neurons
- The marginal probabilities over visible bits are interested
- The hidden bits are the latent representation learned by the network

Training the network



$$P(Y) = \frac{\exp(-E(Y))}{\sum_{s'} \exp(-E(Y'))}$$


$$P(Y) = P(V, H)$$

$$P(V) = \sum_H P(Y)$$

Maximise this term
for training patterns

- $Y=(V, H)$
 - V: output of the visible neurons
 - H: output of the hidden neurons
- The marginal probabilities over visible bits are interested
- The hidden bits are the latent representation learned by the network

Training the network



$$E(Y) = \sum_{i < j} -w_{ij} y_i y_j - b_i y_i$$

$$P(Y) = \frac{\exp(-E(Y))}{\sum_{s'} \exp(-E(Y'))}$$

$$P(V) = \sum_H \frac{\exp(-E(Y))}{\sum_{s'} \exp(-E(Y'))}$$

- Train the network to assign a desired probability distribution to the visible states
- Probability of visible state sums over all hidden states

Maximum likelihood training

$$\log(P(V)) = \log\left(\sum_H \exp\left(\sum_{i<j} w_{ij} y_i y_j\right)\right) - \log\left(\sum_{Y'} \exp\left(\sum_{i<j} w_{ij} y'_i y'_j\right)\right)$$

$$\begin{aligned}\mathcal{L} &= \frac{1}{N} \sum_{V \in \{V\}} \log(P(V)) \\ &= \frac{1}{N} \sum_{V \in \{V\}} \log\left(\sum_H \exp\left(\sum_{i<j} w_{ij} y_i y_j\right)\right) - \log\left(\sum_{Y'} \exp\left(\sum_{i<j} w_{ij} y'_i y'_j\right)\right)\end{aligned}$$

- The loss function is average log likelihood of visible neurons of training vectors $\{V\} = \{V_1, V_2, \dots, V_N\}$
 - should be maximised
 - Two terms have the same format

Maximum likelihood training

$$\mathcal{L} = \frac{1}{N} \sum_{V \in \{V\}} \log \left(\sum_H \exp \left(\sum_{i < j} w_{ij} y_i y_j \right) \right) - \log \left(\sum_{Y'} \exp \left(\sum_{i < j} w_{ij} y'_i y'_j \right) \right)$$
$$\frac{d\mathcal{L}}{dw_{ij}} = \frac{1}{N} \sum_{V \in \{V\}} \sum_H P(Y|V) y_i y_j - \sum_{Y'} P(Y') y'_i y'_j$$

- Similar as the setting without hidden neurons
- But both terms are summations over an exponential states
 - Both need sampling


Maximum likelihood training

$$\begin{aligned}\frac{d\mathcal{L}}{dw_{ij}} &= \frac{1}{N} \sum_{V \in \{V\}} \sum_H P(Y|V) y_i y_j - \sum_{Y'} P(Y') y'_i y'_j \\ \sum_H P(Y|V) y_i y_j &= \frac{1}{K} \sum_{H \in H_{\text{samples}}} y_i y_j \\ \sum_{Y'} P(Y') y'_i y'_j &= \frac{1}{M} \sum_{Y' \in S_{\text{samples}}} y'_i y'_j\end{aligned}$$

- The first term is calculated as the average of sampled hidden state with the visible state fixed
- The second term is calculated as the average of sampled states “freely”


Training Process— Sample 1

- For each training pattern V_i :
 - Fix visible neurons according to V_i
 - Let the hidden neurons evolve from a random initial point to generate H_i
 - Get $Y_i = [V_i, H_i]$
- Repeat K times to generate synthetic training
 - $Y = \{Y_{1,1}, Y_{1,2}, \dots, Y_{1,K}, Y_{2,1}, \dots, Y_{N,K}\}$



Training Process – Sample 2


- Unclamp the visible units and let the entire network evolve several times to generate
 - $Y_{samples} = \{Y_{sample,1}, Y_{sample,2}, \dots, Y_{sample,M}\}$



Training Process

$$\frac{d\mathcal{L}}{dw_{ij}} = \frac{1}{NK} \sum_Y y_i y_j - \frac{1}{M} \sum_{Y' \in Y_{sample}} y'_i y'_j$$
$$w_{ij} = w_{ij} + \alpha \frac{d\mathcal{L}}{dw_{ij}}$$

- Initialise weights
- Get training samples
- Compute gradient and update weights
- Iterate until convergence



Content



- Boltzmann Machine
 - Introduction
 - Training without hidden neurons
 - Training with hidden neurons
 - Summary
- Restricted Boltzmann Machine
- Deep Boltzmann Machine

Boltzmann Machine

- Stochastic extension of Hopfield network
- Store more patterns than Hopfield network through hidden neurons
- Application:
 - Pattern completion
 - Pattern denoising
 - Computing conditional probabilities of patterns
 - Classification
 - Add more bits representing class
 - $[y_1, \dots, y_N, class]$

Boltzmann Machine

- Training process takes a long time...
- Can't work for large problems
- How to solve these problems?


Content



- Boltzmann Machine
 - Introduction
 - Training without hidden neurons
 - Training with hidden neurons
 - Summary
- **Restricted Boltzmann Machine**
- Deep Boltzmann Machine

Restricted Boltzmann machine (RBM)


- Restricted
 - There are no visible-visible and hidden-hidden connections.
 - Proposed as “Harmonium Models” by Paul Smolensky
- Joint Distribution:
 - $$P(V, H) = \frac{\exp(V^T W H + bV + cH)}{\sum_{v', h'} \exp(V'^T W H' + bV' + cH')}$$



Restricted Boltzmann machine (RBM)

$$\text{Hidden: } z_i = \sum_j w_{ji} v_i + b_i \quad P(h_i = 1) = \frac{1}{1 + e^{-z_i}}$$

$$\text{Visible: } y_i = \sum_j w_{ji} h_i + b_i \quad P(v_i = 1) = \frac{1}{1 + e^{-y_i}}$$




- Pros:
 - Sample for hidden neurons: no looping
 - Sample for all neurons: bigraph

Restricted Boltzmann machine (RBM)

$$\text{Hidden: } z_i = \sum_j w_{ji} v_i + b_i \quad P(h_i = 1) = \frac{1}{1 + e^{-z_i}}$$


$$\text{Visible: } y_i = \sum_j w_{ji} h_i + b_i \quad P(v_i = 1) = \frac{1}{1 + e^{-y_i}}$$



- For each sample:
 - Initialize visible neurons
 - Iteratively generate hidden and visible units
 - $\frac{d \log p}{dw_{ij}} = \langle v, h \rangle^0 - \langle v, h \rangle^\infty$

Contrastive Divergence

- Recall in Hopfield Network:
 - No need to raise the entire surface, just the neighborhood
- One iteration is enough in RBM
 - $\frac{d \log p}{dw_{ij}} = \langle v, h \rangle^0 - \langle v, h \rangle^1$



Restricted Boltzmann machine (RBM)

- Generative models for binary data
- Can be extended to continuous-valued data
 - Change the distribution of visible neurons (or hidden neurons)
 - “Exponential Family Harmoniums with an Application to Information Retrieval”, Welling et al., 2004
- Useful for classification and regression

Boltzmann Machines: samples



Content




- Boltzmann Machine
 - Introduction
 - Training without hidden neurons
 - Training with hidden neurons
 - Summary
- Restricted Boltzmann Machine
- **Deep Boltzmann Machine**

Deep Boltzmann Machines

- Stacked RBMs are one of the first deep generative models
- Bottom layer v are visible neurons
- Multiple hidden layers

Deep Boltzmann machine



Boltzmann Machines: samples

Training samples



Generated samples



Reference

- CMU 11-785 Lec 19
- Stanford cs236 Lec 11

Summary

- Boltzmann Machine
 - Introduction
 - Training without hidden neurons
 - Training with hidden neurons
 - Summary
- Restricted Boltzmann Machine **RBM**
- Deep Boltzmann Machine **DBM**

Thanks